



NTNU
Norwegian University of
Science and Technology

Lecture 7: Program Design and Analysis

Asbjørn Djupdal
ARM Norway, IDI NTNU
2013

Lecture overview

- Models of programs
- Compiler toolchain
 - Precompiler
 - Compiler and compilation techniques
 - Assembler
 - Linker
- Optimizing execution time
- Energy aware programming

Models of programs

- Represent the program in a form easy to analyze
- Source code
 - Too hard to manipulate
 - Language specific
- Need a model independent of language
- A compiler uses a temporary representation (IR: Intermediate Representation) for manipulating programs
- Typical data structure: graph

Data flow graph (DFG)

- Models a *basic block*
- A basic block:
 - Code with one start point and one end point where there are no conditional branches
- Gives a partial order of operations in a directed acyclic graph (DAG)
 - Partial because some operations do not depend on each other and can be executed in any order

Single assignment form

- Necessary for creating DFG for a basic block
- Requirement: Each variable is written to only once

$x = a + b;$

$y = c - d;$

$z = x * y;$

$y = b + d;$

original basic block

$x = a + b;$

$y = c - d;$

$z = x * y;$

y1 = b + d;

single assignment form


Data flow graph (DFG)

$x = a + b;$
 $y = c - d;$
 $z = x * y;$
 $y1 = b + d;$

Basic block in single assignment form


Partial orders:

1. $a+b$ and $c-d$
2. $x*y$ and $b+d$



Control data flow graph (CDFG)

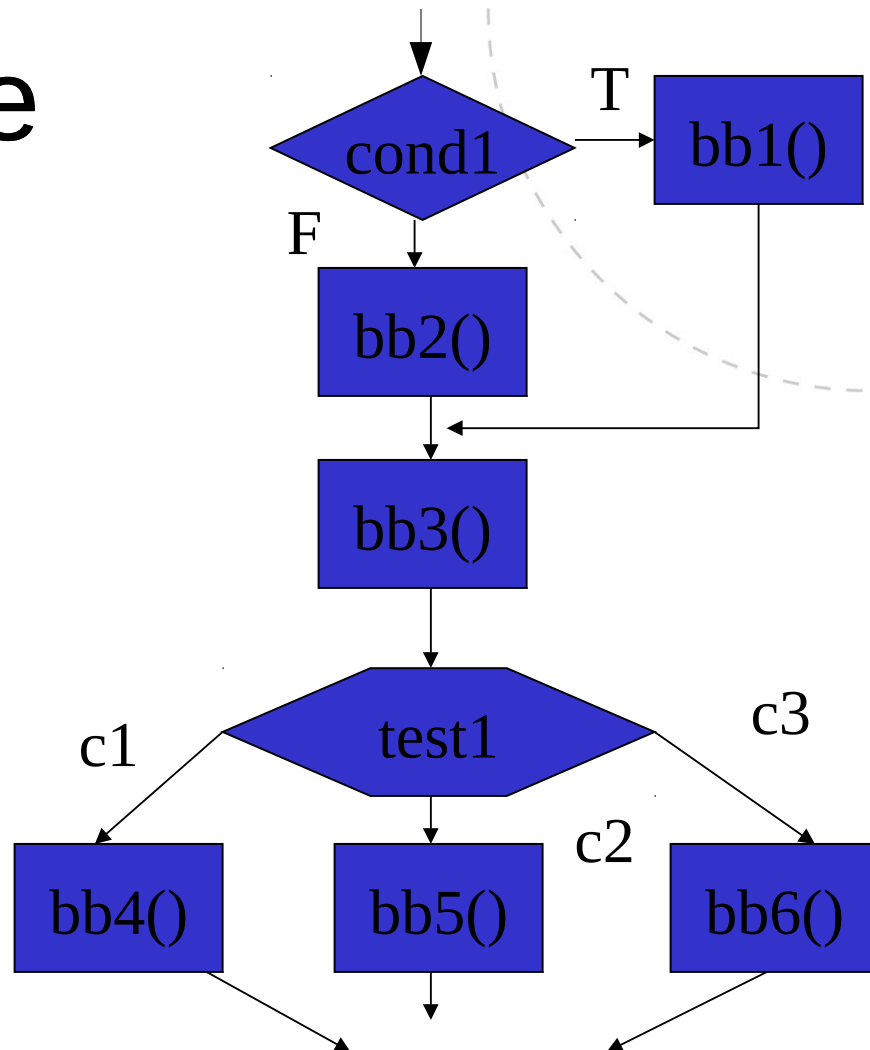
- CDFG: Represents both control and data
- Uses DFG as components
- Two node types:
 - Decision nodes
 - Data flow nodes



basic blocks

CDFG Example

```
if (cond1) {  
    bb1();  
} else {  
    bb2();  
}  
bb3();  
switch (test1) {  
    case c1: bb4(); break;  
    case c2: bb5(); break;  
    case c3: bb6(); break;  
}
```




Loops

for-loop:

```
for (i=0; i<N; i++)  
    loop_body();
```

Equivalent while-loop:

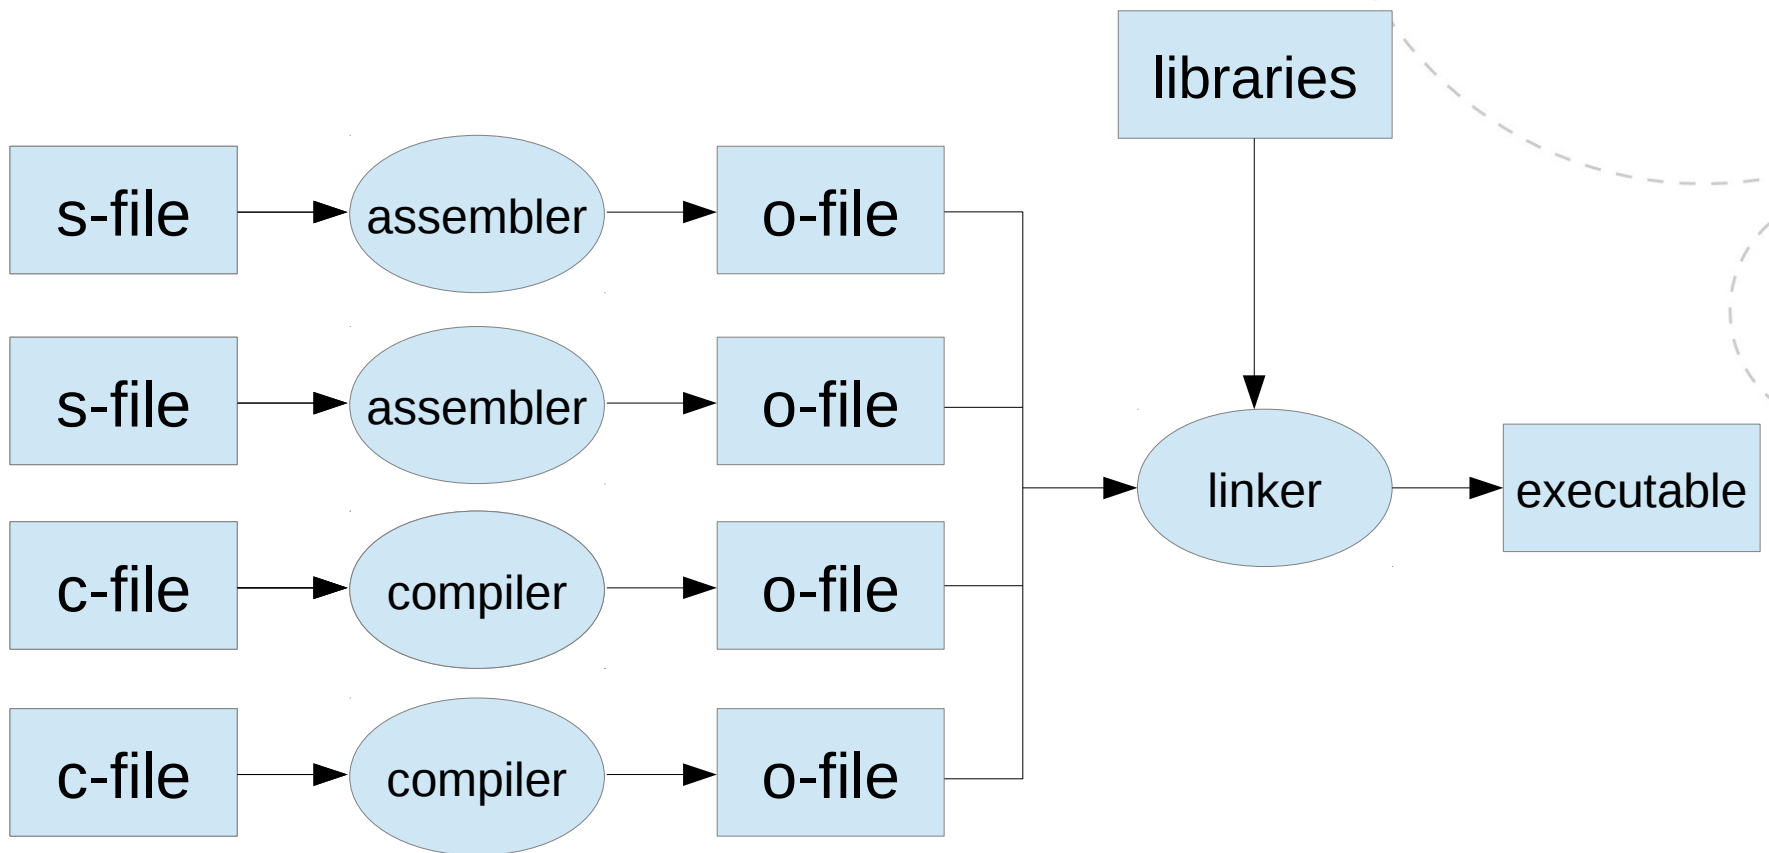
```
i=0;  
while (i<N) {  
    loop_body();  
    i++;  
}
```



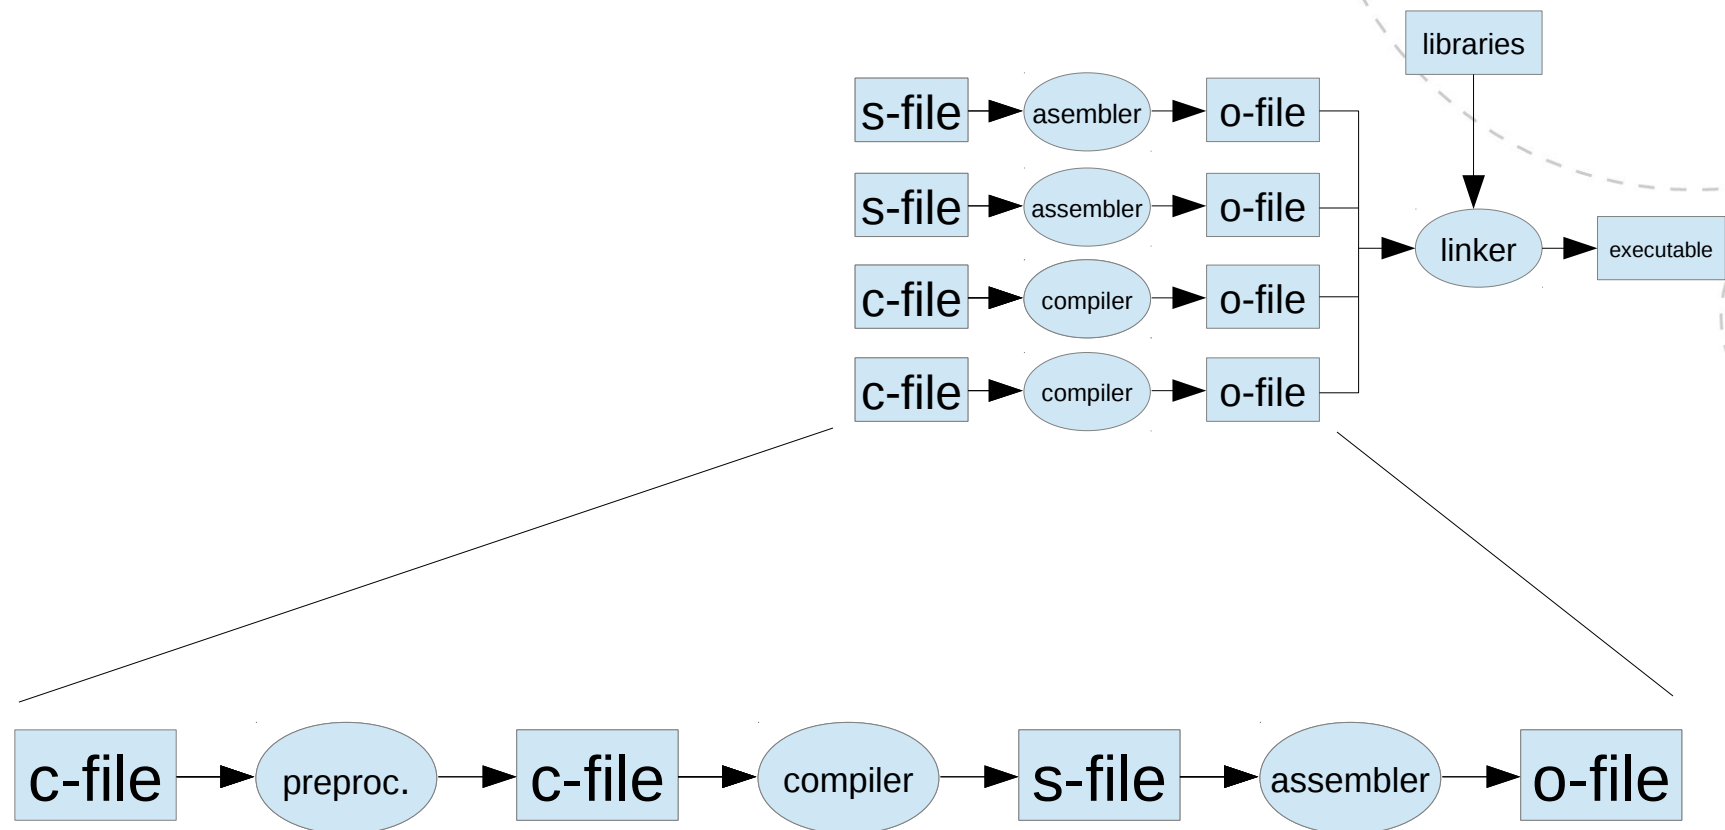
Lecture overview

- Models of programs
- **Compiler toolchain**
 - Precompiler
 - Compiler and compilation techniques
 - Assembler
 - Linker
- Optimizing execution time
- Energy aware programming

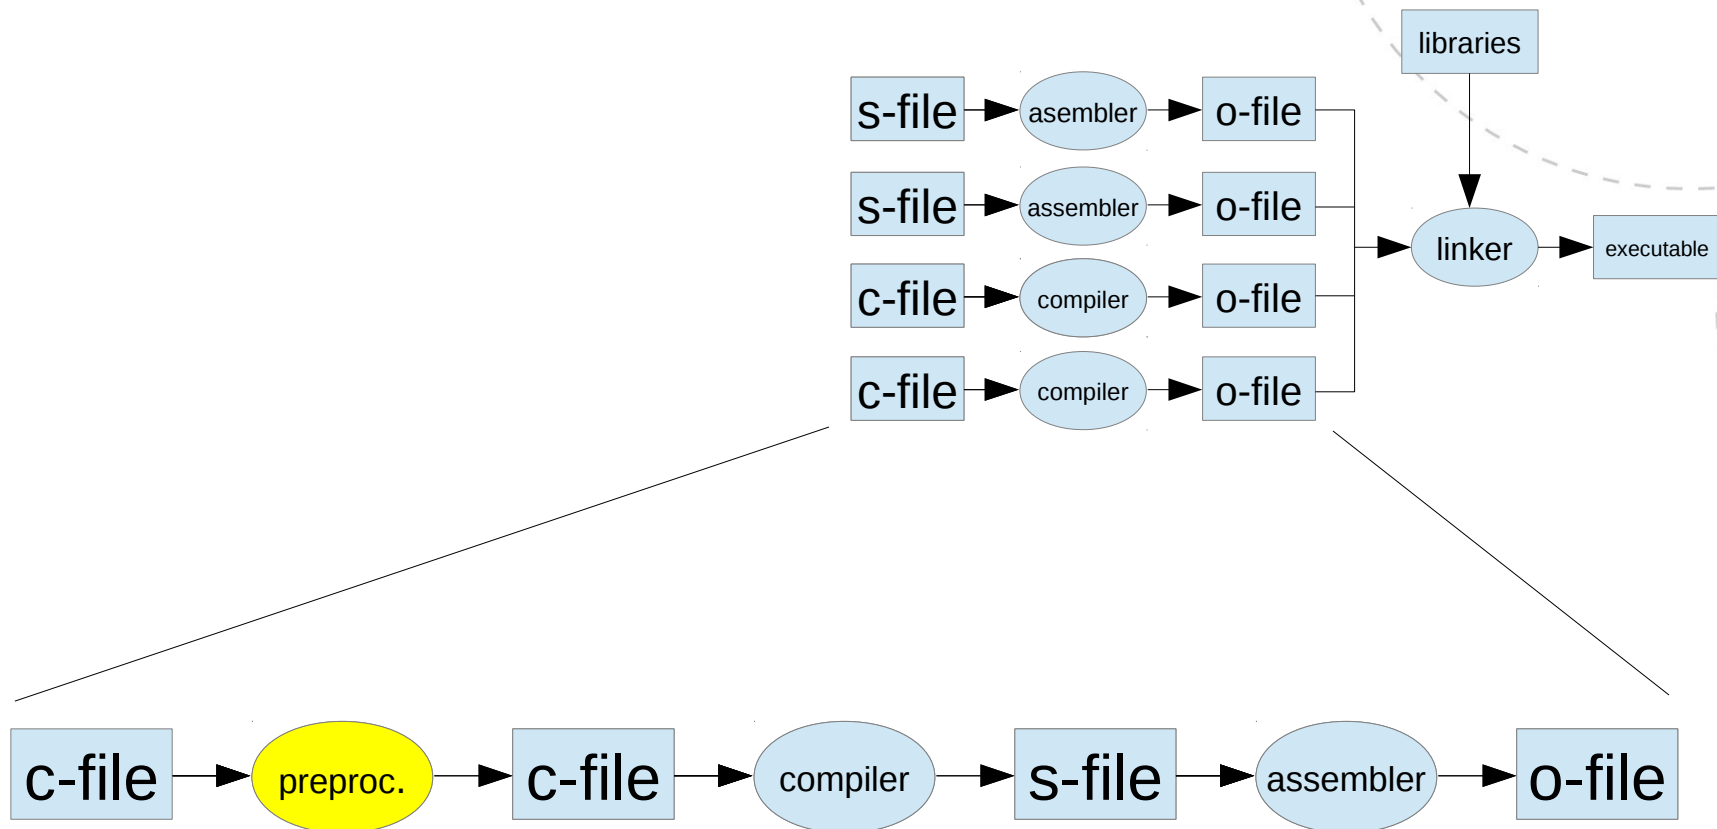
Compiler tool chain



Detailed C tool flow




C preprocessor



C preprocessor

- Transforms C source file with macro substitutions
 - Result is another C source file
- Example: `#define CONST 5`
 - All occurrences of `CONST` in source will be substituted with `5`
 - No type checking or anything, just a direct textual replacement
 - Different from final in java, the following is not equivalent:
 - `#define CONST 5+3`
 - `#define CONST (5+3)`
- Example: `#include "file.h"`
 - The entire `file.h` will be inserted in the C source file
- To examine output from the GCC preprocessor:
 - `gcc -E file.c -o output.c`

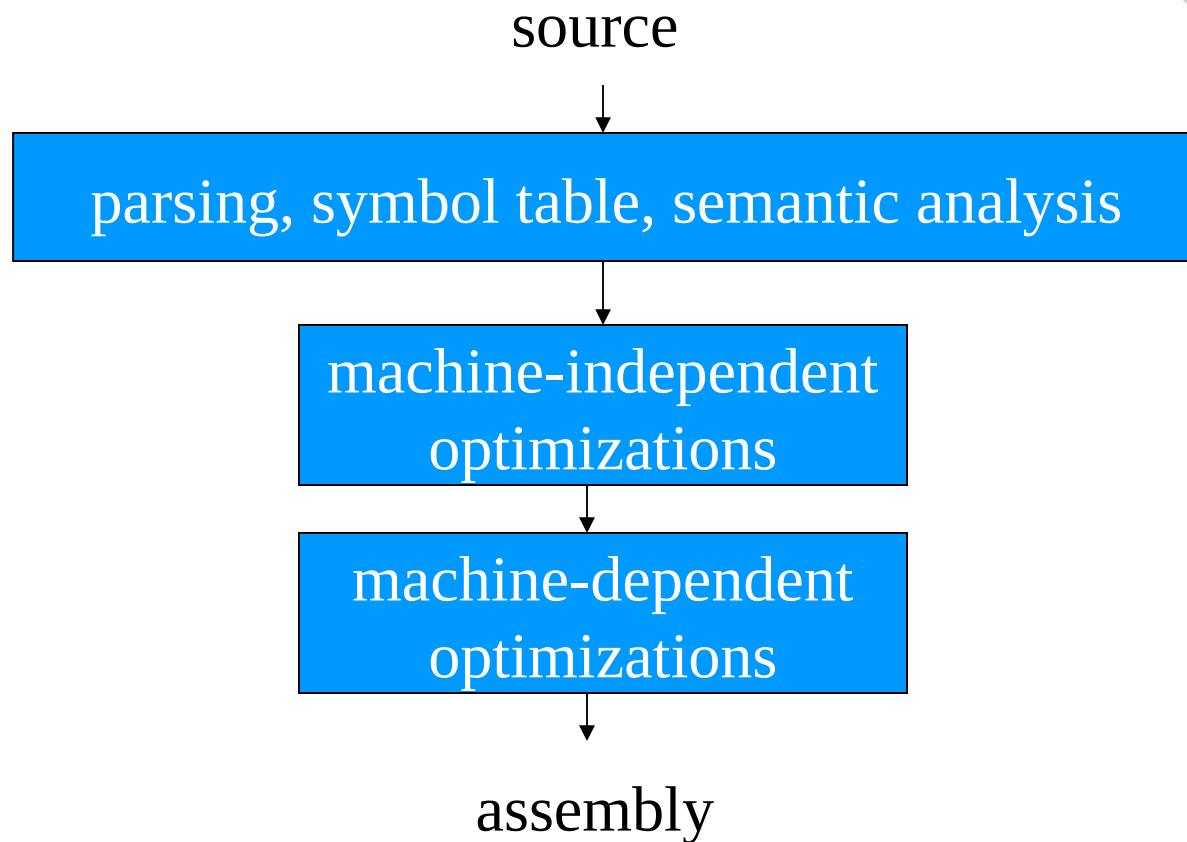
Compilation



Compilation

- Important to understand, especially for embedded systems
 - Intimately related to HW
 - Fewer abstractions (sometimes no OS)
 - Performance, power, energy, real time requirements
 - Which problems can arise during compilations?
 - When is it necessary to write assembly code manually?
- Compiler strategy:
 - Compilation = translation + optimization
- Compiler has a large impact on the resulting program
 - Which and how many instructions are used?
 - How efficient are registers and cache utilized?
 - Are the available resources used?
 - How are memory and I/O operations scheduled?
 - Code size

Compiler phases



Translation and optimization of expressions

- Source code is translated to a temporary CDFG representation
- The CDFG is transformed and optimized
- The CDFG is translated to instructions
- The instructions are further optimized for the target platform

Procedure calls and control flow

- Procedure calls
 - Realized using the processor subroutine call convention
 - Lecture 2 explained this in detail for ARM processors
 - Use the stack to hold local variables and state (saved registers)
 - Pass arguments and return values either through registers or the stack
- Lecture 2 also showed HLL to assembly examples for control flow and expressions
 - Not repeated here

Data structures



- HLL data structures must be realized in memory
 - Integers, floating point, pointers, tables, arrays, trees, graphs, objects, etc
- Typically realized as a base address pointing to a structure with properties known at compile time
- Basic datatypes in C:
 - int, float, double, ...
 - Undefined length
 - Use C99 and include `stdint.h` and `stdbool.h` to get proper datatypes
 - `bool`, `uint32_t`, `int16_t`, ...

Array

- A C array is a pointer to the first element

```
int a[3];
```

a




$= *(a + 1)$

2D array

- Row-major layout in C
- You need to know the size of each row (M) to compute the offset

```
int a[N][M];
```




C struct

- The fields in a struct are static offsets

```
struct mystruct {  
    int field1;  
    char field2;  
};
```

```
struct mystruct a;
```

```
char *aptr = (char*)&a;
```



Warning: alignment


Simplifying expressions

- Constant folding
 - $8 / 2 = 4$
 - `int portaddr = PORT_A + PORT_OFFSET`
- Algebraic
 - $a * b + a * c \longrightarrow a * (b + c)$
- Strength reduction:
 - $a * 2 \longrightarrow a \ll 1;$

Eliminating dead code

- Dead code:

```
#define DEBUG 0  
if(DEBUG) dbg(p1);
```
- Can be eliminated by analyzing the CDFG
 - Reachability analysis
- Reduces code size and may enable additional optimizations



Procedure inlining

- Save procedure call overhead

```
int foo(int a, int b, int c) { return a + b - c; }
```

```
z = foo(w, x, y);
```



```
z = w + x - y;
```

- C99 supports keyword `inline`
 - Can be rejected
 - Optimizer can inline even without keyword
 - But only if the function code is available when compiling
- Can increase code size
- Can reduce cache hit rate

Loop transformations

- A lot of run time is used in loops
- Goals:
 - Reduce loop overhead
 - Increase possibilities for optimizations for pipelines and parallelism
 - Improved use of the memory system

Loop unrolling

- Reduce number of iterations in a loop
- Reduces loop overhead and can enable other optimizations
 - Efficient branch prediction
HW can reduce usefulness unless other optimizations are enabled
 - E.g. reorder instructions from different iterations to eliminate pipeline bubbles

```
for (i=0; i<4; i++)  
    a[i] = b[i] * c[i];
```



```
for (i=0; i<2; i++) {  
    a[i*2] = b[i*2] * c[i*2];  
    a[i*2+1] = b[i*2+1] * c[i*2+1];  
}
```



```
a[0] = b[0]*c[0];  
a[1] = b[1]*c[1];  
a[2] = b[2]*c[2];  
a[3] = b[3]*c[3];
```

Loop fusion

- Combine two loops into one
 - Requires similar iteration space and no dependencies between loops

```
for (i=0; i<N; i++) a[i] = b[i] * 5;  
for (j=0; j<N; j++) w[j] = c[j] * d[j];
```



```
for (i=0; i<N; i++) {  
    a[i] = b[i] * 5;  
    w[i] = c[i] * d[i];  
}
```

Loop fission

- Split one loop into two
 - Cache: Principle of locality
 - (Multicore)

```
for (i=0; i<N; i++) {  
    a[i] = b[i] * 5;  
    w[i] = c[i] * d[i];  
}
```




```
for (i=0; i<N; i++) a[i] = b[i] * 5;  
for (j=0; j<N; j++) w[j] = c[j] * d[j];
```


Loop tiling

- Break a loop into a set of nested loops
 - Each inner loop performs operations on a subset of the data
- Changes the memory access pattern
- Can improve cache behaviour

```
for(i = 0; i < N; i++)
  for(j = 0; j < N; j++)
    f(i,j);
```



```
for(i = 0; i < N; i += 2)
  for(j = 0; j < N; j += 2)
    for(ii = i; ii < min(i + 2, N); i++)
      for(jj = j; jj < min(j + 2, N); j++)
        f(ii,jj);
```



Array padding and struct alignment

- Arrays can be padded with dummy data to better fit cache lines
- Struct elements can be aligned (or padded) to make access more efficient
 - Can also be a requirement by the CPU architecture
- There exists (non-standard) attributes that instructs compiler to avoid aligning and padding

Register allocation


- We need registers to temporarily hold variables
- Choose registers such that memory accesses are minimized
- Analyze the lifetime of variables
- C supports the `register` keyword
 - Hint from the programmer that a variable should be held in a register
 - Usually just ignored by the compiler, rarely used in modern code
- When a C variable *must* be available in memory
 - C keyword `volatile`
 - Still transferred to register, but immediately written back to memory when changed

Registers, lifetime graph

$w = a + b;$ $t=1$


$x = c + w;$ $t=2$

$y = c + d;$ $t=3$



Register allocation

- Method: Conflict graph and graph coloring
 - Edge between variables that are alive at the same time
 - Represent each register with a color
 - Color the nodes with as few colors as possible
 - No edge must share a color
 - NP-complete
 - Compilers use heuristics to find a good solution



3 registers

a r0
b r1
w r2
x r1
c r0
y r2
d r1

Operator scheduling

- Can change the order of instructions to reduce number of registers
 - Changes the life time graph

$w = a + b;$

$x = c + d;$

$y = x + e;$

$z = a - b;$

$w = a + b;$

$z = a - b;$


$x = c + d;$

$y = x + e;$




From 5 registers to 3

Choosing instructions

- A number of instruction sequences can implement a piece of HLL code
- The expressions are represented as a graph
 - Find the best template for the expression
 - Find the mapping that minimizes your chosen cost metric




Expression



Templates

Assembly



Assembly

- Lecture 2 introduced assembly for ARM
- Goals:
 - Generate a binary representation of instructions from a symbolic representation
`ADD R3, R2, R3` \longrightarrow `0xE0823003`
 - Generate symbol table
 - Remember, symbols are:
 - Labels
 - Constants
 - Handle pseudo operations (.word, .data, etc)
- Many assemblers use two passes
 - 1 pass: Generate symbol table
 - 2 pass: Assemble instructions

Symbol table

- Symbols in the assembly file are inserted into the object file *symbol table*.
 - Can then look in the symbol table when assembling instructions referring to symbols
 - Might need to declare in the assembly file which symbols are to be visible to other object files
- Can not always resolve the value of symbols while assembling
 - The start address of the assembly file is often unknown
 - Often necessary to refer to symbols in other object files
- All references to unresolvable symbols are marked as unresolved.
 - Placeholders are inserted where instructions refer to unresolved symbols
 - Must be resolved later, either by the *linker* or the *loader*

Generating the symbol table

CONST = 5

```
        ADD R0, R1, #CONST
label1: ADD R3, R4, R5
        CMP R0, R3
label12: SUB R5, R6, R7
```

Assembly code

CONST = 5

```
label1  = start + 4
label2  = start + 12
```

Symbol Table

Program Location Counter: PLC

The object file

- Result of assembly
- Several standards:
 - ELF (unix), COFF (unix, windows)
- The object file includes:
 - Symbol table
 - Binary program code (text segment)
 - Data (data segment)
 - Uninitialized data (bss segment)
 - Information about relocatable parts
 - Debug data (references to source files)

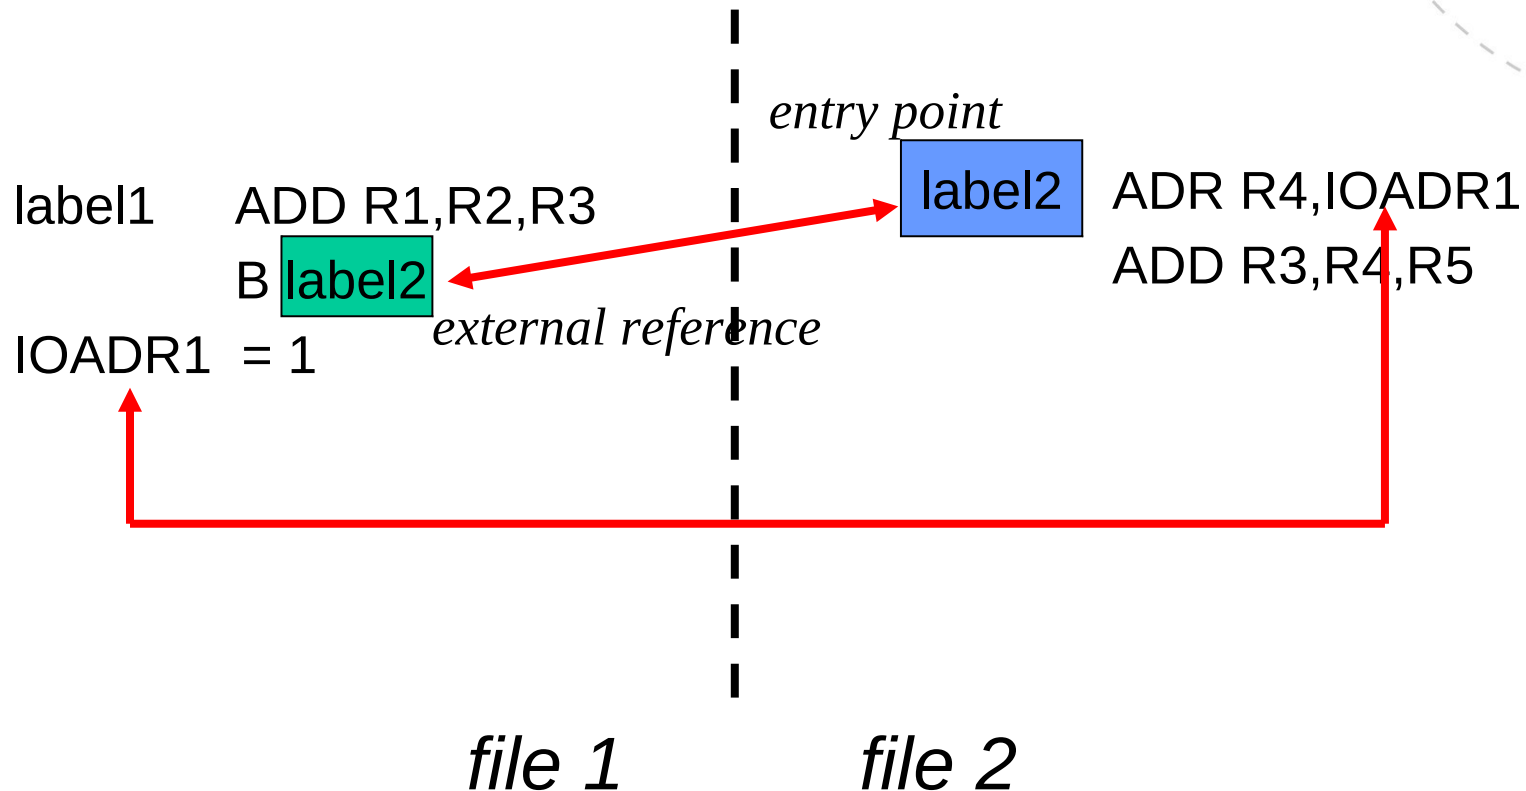
Symbols in C code

- What ends up in the object file symbol table from a C source?

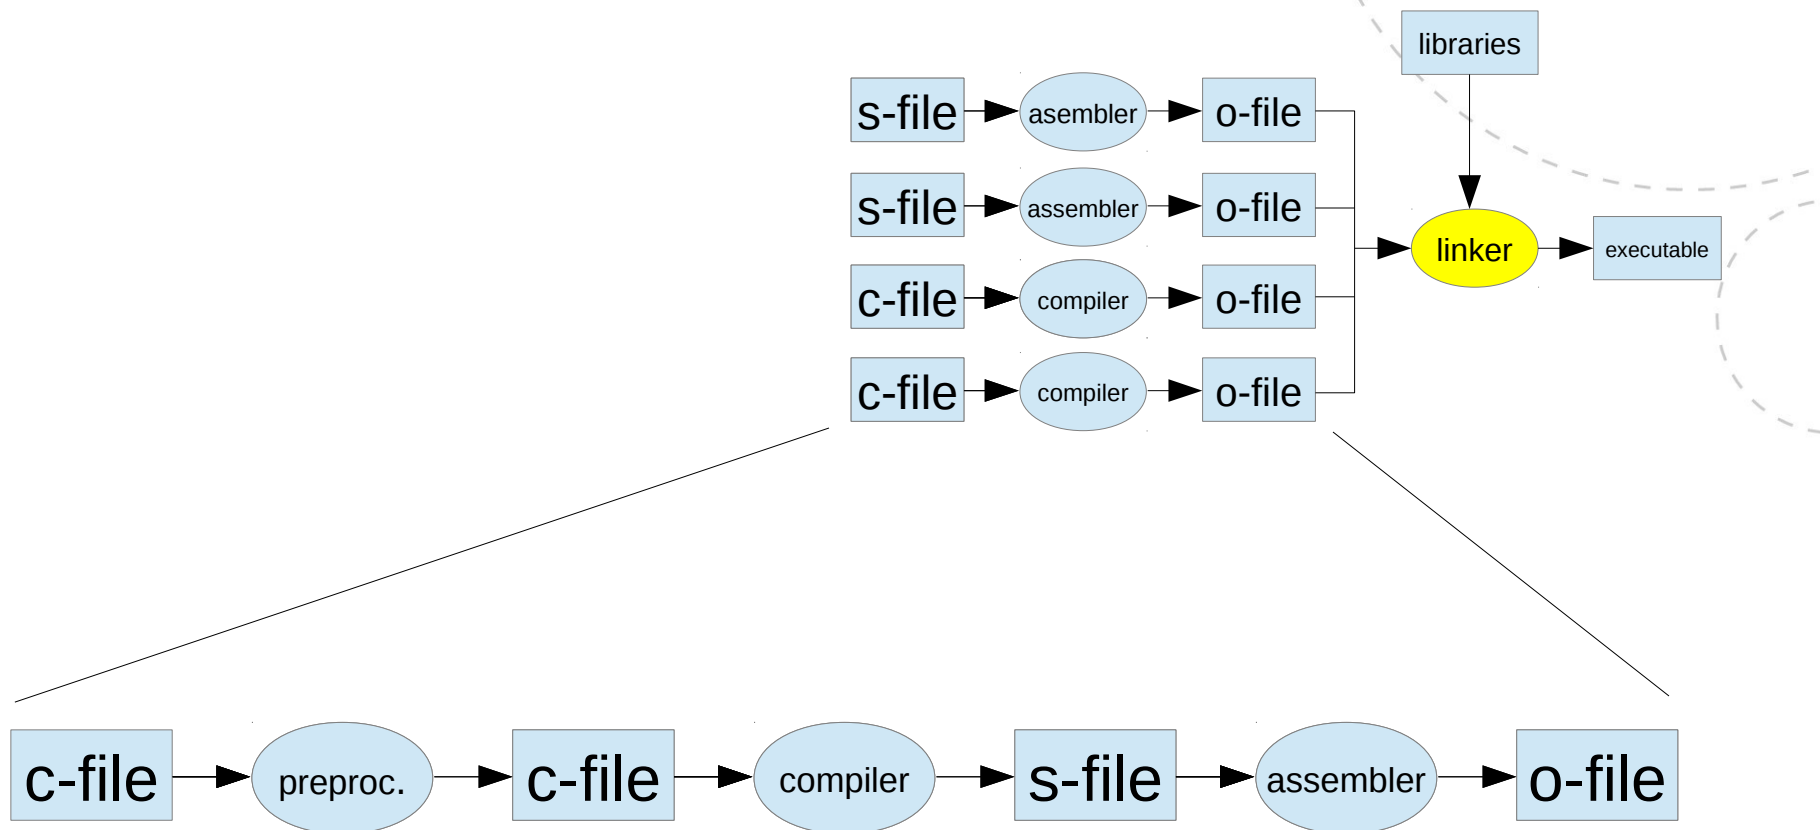
Symbols in C code

- What ends up in the object file symbol table from a C source?
 - Function names
 - But not if declared `static`
 - Global variable names
 - But not if declared `static`
 - Not macros (`#defines`)
- To access symbols from other object files:
 - Need function prototypes to access external functions
 - Need to declare variable `extern` to access external variables
 - `extern int variable;`
 - Typically done in a header file describing the object file external interface

External references and entry points



Linking




Linking

- Takes multiple object files and libraries, and creates one executable file
 - Combines all the object file segments (text, data, bss)
 - Combines all the symbol tables
 - Determines start addresses for all the modules
 - Resolves all symbols
 - Transforms all relative addresses to absolute addresses
 - Fix all unresolved references
 - Can't find a referenced symbol in the merged table? Error!

Module order

- The modules are put in specific locations by the linker
 - Order typically given by the command line argument order
- The linker merges the symbol tables and creates a new symbol table for the merged module
- Addresses in embedded systems often have specific meanings
 - Different addresses can map to different devices/memories (DRAM, EEPROM, cached, non-cached etc)



Custom linker script

- When non-standard behaviour is needed
- Example: Linking for embedded device, without OS
 - Need to tell linker about the stack, heap and which addresses can be used for data and program

```
SECTIONS
{
  . = 0x0;
  .text : {
    startup.o (.text)
    *(.text)
  }
  .data : { *(.data) }
  .bss : { *(.bss COMMON) }
  . = ALIGN(8);
  . = . + 0x1000; /* 4kB of stack memory */
  stack_top = .;
}
```

Dynamic linking

- Most operating systems can link modules at load time
 - Shared libraries
 - .so (linux), .dll (windows)
 - Saves storage space
 - Increases load time
- The compile time linker does not include code from the shared library, only creates references to it
- Small embedded systems do not use shared libraries
- More on dynamic linking and loaders in the next lecture

Interpreters and JIT compilers

- Interpreter
 - Translates and runs expression by expression
 - Slow
 - Can debug, test and run in the same environment
- Just in time-compiler (JIT)
 - Compiles code at runtime
 - Enables different kinds of optimizations
 - Needs warm-up time
 - Stores the translation for use next time
 - Can have large memory requirements

Lecture overview

- Models of programs
- Compiler toolchain
 - Precompiler
 - Compiler and compilation techniques
 - Assembler
 - Linker
- Optimizing execution time
- Energy aware programming

Optimizing execution time

- Embedded systems often react to real time events
 - Need to meet deadlines
 - May need to guarantee that a deadline is always met
- Need to be able to analyze execution time
 - Worst case execution time is often most important
- We need techniques that reduce execution time
- Execution time may vary
 - Input values
 - Cache state
 - Pipeline effects
 - Dynamic interactions of different units in the system

Measuring execution time

- Can be difficult
 - Which input to generate worst case performance?
 - Cache contents
- Software profiler
 - gprof, or similar
 - Changes the code you are profiling
- CPU simulator
 - Different levels of abstraction possible
 - Accuracy vs. simulation speed
 - I/O effects difficult to simulate
- Hardware profiler
 - ARM trace bus
- Logic analyzer

Optimizing execution time

- Make sure you choose the best algorithm
 - No use in optimizing an n^2 algorithm if there exists an $n \cdot \log n$
 - Alg.dat
- Profile and find the parts of the code where time is spent
 - Don't waste your time optimizing if you don't know which parts contribute most to execution time

Reducing program size

- Data footprint
 - Goals
 - Reduce number of main memory accesses
 - Make everything fit inside a small embedded system RAM
 - Reuse constants
 - Reuse variables and data structures
 - If not used at the same time, storage can be reused
- Code size
 - Goals:
 - Reduce instruction cache misses
 - Fit inside limited program memory
 - Choose a processor with compact instructions
 - Use specialized instructions wherever possible
 - Compiler optimizations
 - gcc -Os

Lecture overview

- Models of programs
- Compiler toolchain
 - Precompiler
 - Compiler and compilation techniques
 - Assembler
 - Linker
- Optimizing execution time
- Energy aware programming

Static SW techniques for energy efficiency

- Overall advice: *High performance == low power*
- Execution time optimization
 - Clear correlation between execution time and energy consumption
 - All techniques for reducing execution time will likely reduce energy consumption roughly as much
- Make use of your available HW
 - Unused resources wastes energy
 - Use available resources to make execution time as short as possible
 - Example: If you have support for vector instructions (ARM NEON), their usage can probably help reduce energy consumption
 - Avoid using processors with capabilities you don't need


Static SW techniques for energy efficiency

- Memory access optimization
 - Memory transfers consumes a lot of energy
 - Especially if off-chip
 - Tens to hundreds of times more expensive than an ALU operation
 - Make sure registers are allocated properly to reduce loads and stores
 - Tune algorithms and datastructures such that your working set fits in cache
 - Minimize memory footprint
- I/O access optimization
 - Reduce flash/HD usage
 - Reduce network traffic

Static SW techniques for energy efficiency

- Other static techniques
 - Not very effective
 - Problematic because CPU designers don't give enough information in the datasheet
 - Measurements, trial and error
 - Architecture specific
 - Possible things to try:
 - Some instructions might result in less internal switching than others, even if equally fast
 - Use shift instead of multiplication?
 - Use fixed point (integer instructions) instead of floating point?
 - Instruction sequence can have an impact

Measuring energy consumption



Dynamic energy techniques

- Sleep modes, DVFS, ...
- Has a big and important impact on energy consumption
- Belongs to the OS domain
- Next lecture

Next lecture

- Thursday after easter (?)
 - Operating systems
 - Focus on embedded systems
 - ARM Linux as example
 - Reducing energy consumption with run-time techniques