

Driver development for Linux

Guest lecture TDT4258

2013.04.18

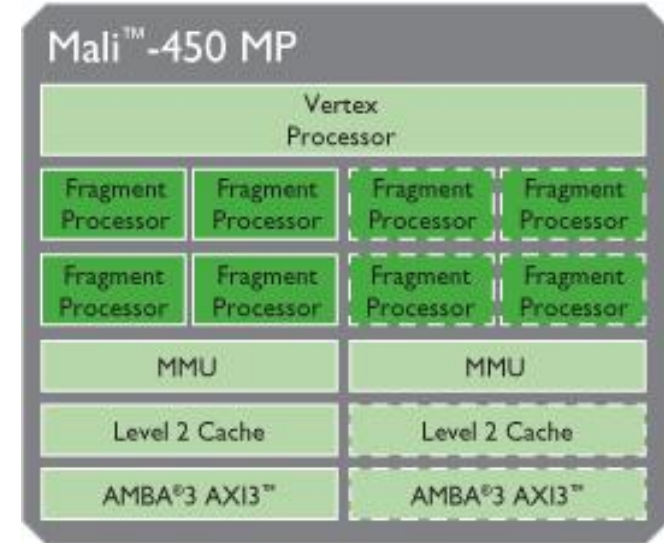
Ørjan Eide and Mikael Valen-Sendstad



The kernel driver for Mali

- Driver for Mali-400

- Mali consist of multiple components
 - Geometry
 - Pixel
 - MMU



- The driver is fairly big

- 25000 code lines kernel driver
- The user space OpenGL driver is bigger, with about half a million code lines

- <http://www.arm.com/products/multimedia/mali-graphics-hardware/index.php>

How the Applications use Mali GPU

- OpenGL ES
 - Subset of OpenGL API designed for use on Embedded Systems
 - Now also for use in web browser through WebGL
- Basic flow:
 - Init: Add textures, vertex buffers, shaders etc
 - Set the render state:
 - Textures
 - Vertices
 - Input: vertex shader (translation matrix)
 - Input: pixel shader (lighting params)
 - Call DrawArray() to draw the object (or DrawElements())
 - Call eglSwapBuffers() to put the buffer on screen

How the Applications use Mali GPU

- Immediate mode rendering
 - Draw calls directly to HW
- Deferred rendering
 - Batch up draw calls
 - Start GP and PP jobs
- Tiled rendering
 - Rendering completes 1 tile at the time
 - Reduces memory bandwidth by caching:
 - Z-read
 - Pixel writes
 - Texture reads

Memory – CPU – GPU

- Computers: Dedicated graphics memory
- Embedded: Shared graphics memory
- Memory latency
- Should graphics memory be cached, by CPU and GPU ?
 - Mapping
 - Cache flushing – Clean, Invalidate
 - Prefetching

Driver architecture

- Most drivers can be built as a kernel module
 - Convenient when developing
 - Allows kernels to support devices without increasing the size of the kernel image
- Linux Driver Model (LDM)
 - Documentation in *Documentation/driver-model/*
 - The model separates the driver from the specification of the device
 - Automatically matches drivers to devices
 - Based on device IDs (for PCI, USB, ...) or a device name
 - Drivers register a table of supported devices
 - Makes it possible to support hot-pluggable devices (e.g. USB)
 - Based on the bus structure in the machine

Partitioning: kernel vs user

- There is often a choice of where to put functionality
 - Whole drivers can run in user space
 - libusb makes it easy to create whole drivers for USB devices in user space
 - It is also possible to create whole applications in kernel space
 - But people rarely do that (for good reasons)
- Robustness and security
 - The consequences of a crash in kernel space can be very big
 - If a user space process crashes it doesn't impact the rest of the system
- Only the kernel can have the full system overview
 - Resource sharing (e.g.: memory or hardware access)
- Performance
 - System calls and context switches costs

Linux driver model

Driver

Device

```
struct platform_device mydevice = {  
    .name = "devicename"  
    .id = 0,  
    .resource = {  
        .start = 0x1000,  
        .end = 0x2000,  
        .name = "device registers",  
        .flags = IORESOURCE_MEM,  
    },  
};  
  
int platform_init_func()  
{  
    return platform_device_register(&mydevice);  
}
```

```
int driver_probe(struct platform_device *pdev)  
{  
    /* Get resources from pdev.  
       Probe the HW, check if the driver is  
       compatible.  
    */  
}  
  
struct platform_driver mydriver = {  
    .probe = &driver_probe,  
    .remove = &driver_remove,  
    .driver = {  
        .name = "devicename",  
        .owner = THIS_MODULE,  
        .bus = &platform_bus_type,  
        .pm = &mydriver_pm_ops,  
    },  
}  
  
int module_init(void)  
{  
    platform_driver_register(&mydriver);  
}
```


Debugging



Debugging

- Reproducing the problem will often be the hardest part
 - Problems that occur once every week, on a machine you can't access across town, is very hard to debug
 - Once you have a good reproducer, fixing the problem is usually relatively easy
- Bugs in kernel space can have non-obvious consequences
 - A simple dangling pointer dereference can cause very unpredictable results
- Without visibility of what the driver is doing it is hard to debug
 - There are many ways to gain some visibility
 - Log actions as they happen
 - Could be done simply with **printk**, or with more advanced logging
 - Use a debugger: kgdb, JTAG
 - Pre-emptive debugging built into the driver or kernel
 - Extra checks, asserts, ...

Typical debugging workflow

- 1. Reproduce problem**
- 2. Gather information**
- 3. Try to understand problem**
 - 1. If more information is needed:**
 - 1. Increase visibility of what is happening**
 - 2. Goto 1.**
- 4. Fix problem**
- 5. Verify fix**
- 6. Create test**
 - Catches error if reintroduced

Debugging: Debuggers

- Hardware debuggers: JTAG
 - It can be hard to gain visibility of software through software
 - The debugging/tracing can interfere with the driver
 - Quite often the problem disappears as soon as you attach gdb
 - A HW debugger lets you stop the CPU and inspect memory and registers
 - Can also debug interrupt handlers and regions of code running with interrupts disabled
- gdb for the Linux kernel: kgdb
 - Connects to a running kernel over serial port or network
 - Can't do everything you can do with a HW debugger
 - But, you don't need specialized (possibly expensive) hardware

Debugging: Linux

- There are a lot of great debugging aids built into the Linux kernel
- Check «Kernel hacking» in menuconfig
 - Many of these functions helps detection problems early
- Using a kernel with most of these debugging aids enabled can be good during development
 - It will be slower than a "normal" kernel
 - But it is worth it to discover problems early
 - Some of the debugging aids have only a very small performance impact and can be enabled even for production kernel builds

Debugging: Linux

- **kmemleak**
 - Tracks all memory allocations and discovers memory leaks
 - *Documentation/kmemleak.txt*
- **The built-in Linux lock debugging is great**
 - It can be quite hard to do locking correctly
 - Errors can lead to hanging processes, or whole systems
 - **CONFIG_DEBUG_MUTEXES, CONFIG_DEBUG_SPINLOCK**
 - Does some basic checks for locks
 - **CONFIG_PROVE_LOCKING**
 - Proves that all observed locking behavior is safe
 - Generates a warning if a deadlock could possibly occur
 - **And many more:** CONFIG_DEBUG_ATOMIC_SLEEP, CONFIG_DEBUG_RT_MUTEXES, CONFIG_DEBUG_LOCK_ALLOC, ...

Debugging: Linux

■ Ftrace/Tracepoints

- *Documentation/trace/ftrace.txt*
- *Documentation/trace/tracepoints.txt*

- Traces events and function calls in the kernel
- A driver can emit events for its own events

```
$ cd /sys/kernel/debug/tracing
$ echo 1 > /sys/kernel/debug/tracing/tracing_enabled
$ cat /sys/kernel/debug/tracing/trace > /tmp/trace.txt
$ echo 0 > /sys/kernel/debug/tracing/tracing_enabled
```

```
# tracer: function_graph
#
# TIME CPU DURATION FUNCTION CALLS
# | | | | |
3) 0.105 us | } /* rcu_enter_nohz */
3) 1.083 us | } /* rcu_irq_exit */
3) 0.066 us | idle_cpu();
3) | tick_nohz_stop_sched_tick() {
3) | ktime_get() {
3) 1.125 us | read_hpet();
3) 1.626 us | }
3) 0.066 us | update_ts_time_stats();
3) 0.066 us | timekeeping_max_deferment();
3) 0.171 us | rcu_needs_cpu();
3) 0.069 us | printk_needs_cpu();
3) | get_next_timer_interrupt() {
3) 0.096 us | _raw_spin_lock();
3) | hrtimer_get_next_event() {
3) 0.105 us | _raw_spin_lock_irqsave();
3) 0.078 us | _raw_spin_unlock_irqrestore();
3) 1.095 us | }
3) 2.181 us | }
3) | hrtimer_start() {
3) | __hrtimer_start_range_ns() {
3) | lock_hrtimer_base() {
3) 0.099 us | _raw_spin_lock_irqsave();
3) 0.621 us | }
3) 0.114 us | __remove_hrtimer();
3) 0.201 us | enqueue_hrtimer();
3) 0.078 us | _raw_spin_unlock_irqrestore();
3) 2.841 us | }
3) 3.417 us | }
3) + 10.893 us | }
3) ! 110.298 us | } /* irq_exit */
3) ! 125.199 us | } /* smp_call_function_single_interrupt */
```

ARM DS-5

- IDE with built in debugging and trace features
 - Supports HW debugger, DSTREAM
 - And remote debug similar to kgdb



- Streamline
 - Remote debug and trace
 - Trace based on Linux ftrace and tracepoints
 - Uses the Linux perf framework to gather CPU performance counters

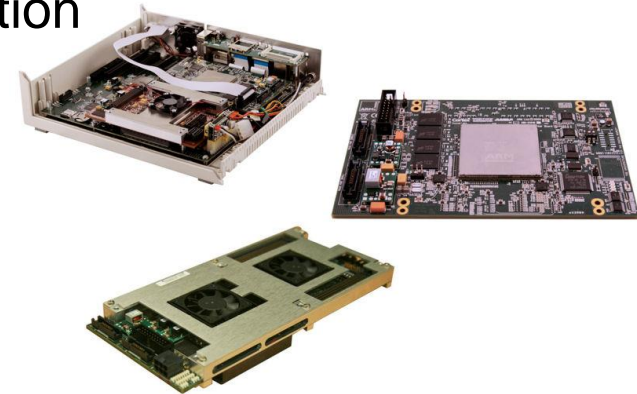


- <http://www.arm.com/products/tools/software-tools/ds-5/index.php>

Driver development without hardware

■ Prototype in FPGA

- Can be used to develop driver before HW is ready
- Driver development can feed back into HW development
- Probably slower than real HW
- Software can contribute to hardware verification



■ Software model

- E.g. ARM Fast Models
 - <http://www.arm.com/products/tools/models/fast-models/index.php>
- Can start driver development before even FPGA is available

Questions?

